

Find minimum of unconstrained multivariable function

Equation

Finds the minimum of a problem specified by

$$\min_x f(x)$$

where $f(x)$ is a function that returns a scalar.

x is a vector or a matrix; see [Matrix Arguments](#).

Syntax

```
x = fminunc(fun,x0)
x = fminunc(fun,x0,options)
x = fminunc(problem)
[x,fval] = fminunc(...)
[x,fval,exitflag] = fminunc(...)
[x,fval,exitflag,output] = fminunc(...)
[x,fval,exitflag,output,grad] = fminunc(...)
[x,fval,exitflag,output,grad,hessian] = fminunc(...)
```

Description

fminunc attempts to find a minimum of a scalar function of several variables, starting at an initial estimate. This is generally referred to as *unconstrained nonlinear optimization*.

Note: [Passing Extra Parameters](#) explains how to pass extra parameters to the objective function, if necessary.

`x = fminunc(fun,x0)` starts at the point `x0` and attempts to find a local minimum `x` of the function described in `fun`. `x0` can be a scalar, vector, or matrix.

`x = fminunc(fun,x0,options)` minimizes with the optimization options specified in `options`. Use [optimoptions](#) to set these options.

`x = fminunc(problem)` finds the minimum for `problem`, where `problem` is a structure described in [Input Arguments](#).

Create the problem structure by exporting a problem from Optimization app, as described in [Exporting Your Work](#).

`[x,fval] = fminunc(...)` returns in `fval` the value of the objective function `fun` at the solution `x`.

`[x,fval,exitflag] = fminunc(...)` returns a value `exitflag` that describes the exit condition.

`[x,fval,exitflag,output] = fminunc(...)` returns a structure `output` that contains information about the optimization.

`[x,fval,exitflag,output,grad] = fminunc(...)` returns in `grad` the value of the gradient of `fun` at the solution `x`.

`[x,fval,exitflag,output,grad,hessian] = fminunc(...)` returns in `hessian` the value of the Hessian of the objective function `fun` at the solution `x`. See [Hessian](#).

Input Arguments

[Function Arguments](#) contains general descriptions of arguments passed into `fminunc`. This section provides function-specific details for `fun`, `options`, and `problem`:

fun	<p>The function to be minimized. <code>fun</code> is a function that accepts a vector <code>x</code> and returns a scalar <code>f</code>, the objective function evaluated at <code>x</code>. The function <code>fun</code> can be specified as a function handle for a file</p> <pre>x = fminunc(@myfun,x0)</pre> <p>where <code>myfun</code> is a MATLAB[®] function such as</p> <pre>function f = myfun(x)</pre>
-----	--

	<pre>f = ... % Compute function value at x</pre> <p>fun can also be a function handle for an anonymous function.</p> <pre>x = fminunc(@(x)norm(x)^2,x0);</pre> <p>If the gradient of fun can also be computed <i>and</i> the GradObj option is 'on', as set by</p> <pre>options = optimoptions(@fminunc,'GradObj','on')</pre> <p>then the function fun must return, in the second output argument, the gradient value g, a vector, at x. The gradient is the partial derivatives $\partial f/\partial x_i$ of f at the point x. That is, the <i>i</i>th component of g is the partial derivative of f with respect to the <i>i</i>th component of x.</p> <p>If the Hessian matrix can also be computed <i>and</i> the Hessian option is 'on', i.e., options = optimoptions(@fminunc,'GradObj','on','Hessian','on'), then the function fun must return the Hessian value H, a symmetric matrix, at x in a third output argument. The Hessian matrix is the second partial derivatives matrix of f at the point x. That is, the (i,j)th component of H is the second partial derivative of f with respect to x_i and x_j, $\partial^2 f/\partial x_i \partial x_j$. The Hessian is by definition a symmetric matrix.</p> <p>Writing Scalar Objective Functions explains how to "conditionalize" the gradients and Hessians for use in solvers that do not accept them. Passing Extra Parameters explains how to parameterize fun, if necessary.</p>	
options	Options provides the function-specific details for the options values.	
problem	objective	Objective function
	x0	Initial point for x
	solver	'fminunc'
	options	Options created with optimoptions

Output Arguments

[Function Arguments](#) contains general descriptions of arguments returned by fminunc. This section provides function-specific details for exitflag and output:

exitflag	Integer identifying the reason the algorithm terminated. The following lists the values of exitflag and the corresponding reasons the algorithm terminated.	
	1	Magnitude of gradient smaller than the TolFun tolerance.
	2	Change in x was smaller than the TolX tolerance.
	3	Change in the objective function value was less than the TolFun tolerance.
	5	Predicted decrease in the objective function was less than the TolFun tolerance.
	0	Number of iterations exceeded MaxIter or number of function evaluations exceeded MaxFunEvals.
	-1	Algorithm was terminated by the output function.
	-3	Objective function at current iteration went below ObjectiveLimit.
grad	Gradient at x	
hessian	Hessian at x	
output	Structure containing information about the optimization. The fields of the structure are	
	iterations	Number of iterations taken

funcCount	Number of function evaluations
firstorderopt	Measure of first-order optimality
algorithm	Optimization algorithm used
cgiterations	Total number of PCG iterations ('trust-region ' algorithm only)
stepsize	Final displacement in x ('quasi-newton ' algorithm only)
message	Exit message

Hessian

fminunc computes the output argument hessian as follows:

- When using the 'quasi-newton' algorithm, the function computes a finite-difference approximation to the Hessian at x using
 - The gradient grad if you supply it
 - The objective function fun if you do not supply the gradient
- When using the 'trust-region' algorithm, the function uses
 - options.Hessian, if you supply it, to compute the Hessian at x
 - A finite-difference approximation to the Hessian at x, if you supply only the gradient

Options

fminunc uses these optimization options. Some options apply to all algorithms, some are only relevant when you are using the trust-region algorithm, and others are only relevant when you are using the quasi-newton algorithm. Use [optimoptions](#) to set or change options. See [Optimization Options Reference](#) for detailed information.

All Algorithms

All fminunc algorithms use the following options:

Algorithm If you use optimset (not recommended), use LargeScale instead of Algorithm.	Choose the fminunc algorithm. Choices are 'quasi-newton' and 'trust-region' (default). The trust-region algorithm requires you to provide the gradient (see the preceding description of fun), or else fminunc uses the 'quasi-newton' algorithm. For information on choosing the algorithm, see Choosing the Algorithm .
DerivativeCheck	Compare user-supplied derivatives (gradient of objective) to finite-differencing derivatives. The choices are 'on' or the default 'off'.
Diagnostics	Display diagnostic information about the function to be minimized or solved. The choices are 'on' or the default 'off'.
DiffMaxChange	Maximum change in variables for finite-difference gradients (a positive scalar). The default is Inf.
DiffMinChange	Minimum change in variables for finite-difference gradients (a positive scalar). The default is 0.
Display	Level of display: <ul style="list-style-type: none"> 'off' or 'none' displays no output. 'iter' displays output at each iteration, and gives the default exit message. 'iter-detailed' displays output at each iteration, and gives the technical exit message. 'notify' displays output only if the function does not converge, and gives the default exit message. 'notify-detailed' displays output only if the function does not converge, and gives the technical exit message.

	<ul style="list-style-type: none"> 'final' (default) displays just the final output, and gives the default exit message. 'final-detailed' displays just the final output, and gives the technical exit message.
FinDiffRelStep	<p>Scalar or vector step size factor. When you set FinDiffRelStep to a vector v, forward finite differences δ are</p> $\delta = v.\text{sign}(x).\text{max}(\text{abs}(x), \text{TypicalX});$ <p>and central finite differences are</p> $\delta = v.\text{max}(\text{abs}(x), \text{TypicalX});$ <p>Scalar FinDiffRelStep expands to a vector. The default is $\sqrt{\text{eps}}$ for forward finite differences, and $\text{eps}^{(1/3)}$ for central finite differences.</p> <p>The trust-region algorithm uses FinDiffRelStep only when DerivativeCheck is 'on'.</p>
FinDiffType	<p>Finite differences, used to estimate gradients, are either 'forward' (the default), or 'central' (centered). 'central' takes twice as many function evaluations, but should be more accurate. The trust-region algorithm uses FinDiffType only when DerivativeCheck is 'on'.</p>
FunValCheck	<p>Check whether objective function values are valid. 'on' displays an error when the objective function returns a value that is complex, Inf, or NaN. The default, 'off', displays no error.</p>
GradObj	<p>Gradient for the objective function defined by the user. See the preceding description of fun to see how to define the gradient in fun. Set to 'on' to have fminunc use a user-defined gradient of the objective function. The default 'off' causes fminunc to estimate gradients using finite differences. You must provide the gradient, and set GradObj to 'on', to use the trust-region algorithm. This option is not required for the quasi-Newton algorithm.</p>
LargeScale	<p>Use 'trust-region' algorithm if possible when set to the default 'on'. Use 'quasi-newton' algorithm when set to 'off'.</p> <p>If you use optimoptions (recommended), use Algorithm instead of LargeScale.</p> <p>The LargeScale algorithm requires you to provide the gradient (see the preceding description of fun), or else fminunc uses the 'quasi-newton' algorithm. For information on choosing the algorithm, see Choosing the Algorithm.</p>
MaxFunEvals	<p>Maximum number of function evaluations allowed, a positive integer. The default value is 100*numberOfVariables.</p>
MaxIter	<p>Maximum number of iterations allowed, a positive integer. The default value is 400.</p>
OutputFcn	<p>Specify one or more user-defined functions that an optimization function calls at each iteration, either as a function handle or as a cell array of function handles. The default is none ([]). See Output Function.</p>
PlotFcns	<p>Plots various measures of progress while the algorithm executes. Select from predefined plots or write your own. Pass a function handle or a cell array of function handles. The default is none ([]).</p> <ul style="list-style-type: none"> @optimplotx plots the current point. @optimplotfunccount plots the function count. @optimplotfval plots the function value. @optimplotstepsize plots the step size. @optimplotfirstorderopt plots the first-order optimality measure. <p>For information on writing a custom plot function, see Plot Functions.</p>
TolFun	<p>Termination tolerance on the function value, a positive scalar. The default is 1e-6.</p>
TolX	<p>Termination tolerance on x, a positive scalar. The default value is 1e-6.</p>

TypicalX	<p>Typical x values. The number of elements in TypicalX is equal to the number of elements in x0, the starting point. The default value is ones(numberofvariables,1). fminunc uses TypicalX for scaling finite differences for gradient estimation.</p> <p>The trust-region algorithm uses TypicalX only for the DerivativeCheck option.</p>
----------	--

trust-region Algorithm Only

The trust-region algorithm uses the following options:

Hessian	<p>If 'on', fminunc uses a user-defined Hessian for the objective function. The Hessian is either defined in the objective function (see fun), or is indirectly defined when using HessMult.</p> <p>If 'off' (default), fminunc approximates the Hessian using finite differences.</p>
HessMult	<p>Function handle for Hessian multiply function. For large-scale structured problems, this function computes the Hessian matrix product H*Y without actually forming H. The function is of the form</p> $W = \text{hmfun}(\text{Hinfo}, Y)$ <p>where Hinfo contains the matrix used to compute H*Y.</p> <p>The first argument must be the same as the third argument returned by the objective function fun, for example by</p> $[f, g, \text{Hinfo}] = \text{fun}(x)$ <p>Y is a matrix that has the same number of rows as there are dimensions in the problem. $W = H*Y$ although H is not formed explicitly. fminunc uses Hinfo to compute the preconditioner. See Passing Extra Parameters for information on how to supply values for any additional parameters hmfun needs.</p> <div style="border: 1px solid black; padding: 5px; margin-top: 10px;"> <p>Note 'Hessian' must be set to 'on' for fminunc to pass Hinfo from fun to hmfun.</p> </div> <p>See Minimization with Dense Structured Hessian, Linear Equalities for an example.</p>
HessPattern	<p>Sparsity pattern of the Hessian for finite differencing. Set HessPattern(i,j) = 1 when you can have $\partial^2 \text{fun} / \partial x(i) \partial x(j) \neq 0$. Otherwise, set HessPattern(i,j) = 0.</p> <p>Use HessPattern when it is inconvenient to compute the Hessian matrix H in fun, but you can determine (say, by inspection) when the ith component of the gradient of fun depends on x(j). fminunc can approximate H via sparse finite differences (of the gradient) if you provide the <i>sparsity structure</i> of H — i.e., locations of the nonzeros — as the value for HessPattern.</p> <p>In the worst case, when the structure is unknown, do not set HessPattern. The default behavior is as if HessPattern is a dense matrix of ones. Then fminunc computes a full finite-difference approximation in each iteration. This can be very expensive for large problems, so it is usually better to determine the sparsity structure.</p>
MaxPCGIter	Maximum number of PCG (preconditioned conjugate gradient) iterations, a positive scalar. The default is max(1, floor(numberOfVariables/2)). For more information, see Algorithms .
PrecondBandWidth	Upper bandwidth of preconditioner for PCG, a nonnegative integer. By default, fminunc uses diagonal preconditioning (upper bandwidth of 0). For some problems, increasing the bandwidth reduces the number of PCG iterations. Setting PrecondBandWidth to Inf uses a direct factorization (Cholesky) rather than the conjugate gradients (CG). The direct factorization is computationally more expensive than CG, but produces a better quality step towards the solution.
TolPCG	Termination tolerance on the PCG iteration, a positive scalar. The default is 0.1.

quasi-newton Algorithm Only

The quasi-newton algorithm uses the following options:

HessUpdate	Method for choosing the search direction in the Quasi-Newton algorithm. The choices are: <ul style="list-style-type: none"> 'bfgs', the default 'dfp' 'steepdesc' See Hessian Update for a description of these methods.
InitialHessMatrix This option will be removed in a future release.	Initial quasi-Newton matrix. This option is only available if you set InitialHessType to 'user-supplied'. In that case, you can set InitialHessMatrix to one of the following: <ul style="list-style-type: none"> A positive scalar — The initial matrix is the scalar times the identity. A vector of positive values— The initial matrix is a diagonal matrix with the entries of the vector on the diagonal. This vector should be the same size as the x0 vector, the initial point.
InitialHessType This option will be removed in a future release.	Initial quasi-Newton matrix type. The options are: <ul style="list-style-type: none"> 'identity' 'scaled-identity', the default 'user-supplied' — See InitialHessMatrix
ObjectiveLimit	A tolerance (stopping criterion) that is a scalar. If the objective function value at an iteration is less than or equal to ObjectiveLimit, the iterations halt, since the problem is presumably unbounded. The default value is -1e20.

Examples

Minimize the function $f(x) = 3x_1^2 + 2x_1x_2 + x_2^2$.

Create a file myfun.m:

```
function f = myfun(x)
f = 3*x(1)^2 + 2*x(1)*x(2) + x(2)^2;    % Cost function
```

Then call fminunc to find a minimum of myfun near [1,1]:

```
x0 = [1,1];
[x,fval] = fminunc(@myfun,x0);
```

After a few iterations, fminunc returns the solution, x, and the value of the function at x, fval:

```
x,fval

x =
    1.0e-006 *
    0.2541    -0.2029

fval =
    1.3173e-013
```

To minimize this function with the gradient provided, modify myfun.m so the gradient is the second output argument:

```
function [f,g] = myfun(x)
f = 3*x(1)^2 + 2*x(1)*x(2) + x(2)^2;    % Cost function
if nargin > 1
    g(1) = 6*x(1)+2*x(2);
    g(2) = 2*x(1)+2*x(2);
end
```

Indicate that the gradient value is available by creating optimization options with the GradObj option set to 'on' using [optimoptions](#). Choose the 'trust-region' algorithm, which requires a gradient.

```
options = optimoptions('fminunc','GradObj','on','Algorithm','trust-region');
x0 = [1,1];
[x,fval] = fminunc(@myfun,x0,options);
```

After several iterations fminunc returns the solution, x, and the value of the function at x, fval:

```
x,fval

x =
    1.0e-015 *
    0.1110    -0.8882

fval =
    6.2862e-031
```

To minimize the function $f(x) = \sin(x) + 3$ using an anonymous function

```
f = @(x)sin(x)+3;
x = fminunc(f,4);
```

fminunc returns a solution

```
x

x =
    4.7124
```

Notes

fminunc is not the preferred choice for solving problems that are sums of squares, that is, of the form

$$\min_x \|f(x)\|_2^2 = \min_x \left(f_1(x)^2 + f_2(x)^2 + \dots + f_n(x)^2 \right)$$

Instead use the [lsqnonlin](#) function, which has been optimized for problems of this form.

To use the trust-region method, you must provide the gradient in fun (and set the GradObj option to 'on' using optimoptions). A warning is given if no gradient is provided and the Algorithm option is 'trust-region'.

Limitations

The function to be minimized must be continuous. fminunc might only give local solutions.

fminunc only minimizes over the real numbers, that is, x must only consist of real numbers and f(x) must only return real numbers. When x has complex variables, they must be split into real and imaginary parts.

Trust Region Optimization

To use the trust-region algorithm, you must supply the gradient in fun (and GradObj must be set 'on' in options).

Trust Region Algorithm Coverage and Requirements

Additional Information Needed	For Large Problems
Must provide gradient for f(x) in fun.	<ul style="list-style-type: none">Provide sparsity structure of the Hessian, or compute the Hessian in fun.The Hessian should be sparse.

More About

[expand all](#)

Algorithms

- [function_handle](#)
- [Anonymous Functions](#)

- [Unconstrained Optimization](#)

References

- [1] Broyden, C.G., "The Convergence of a Class of Double-Rank Minimization Algorithms," *Journal Inst. Math. Applic.*, Vol. 6, pp. 76-90, 1970.
- [2] Coleman, T.F. and Y. Li, "An Interior, Trust Region Approach for Nonlinear Minimization Subject to Bounds," *SIAM Journal on Optimization*, Vol. 6, pp. 418-445, 1996.
- [3] Coleman, T.F. and Y. Li, "On the Convergence of Reflective Newton Methods for Large-Scale Nonlinear Minimization Subject to Bounds," *Mathematical Programming*, Vol. 67, Number 2, pp. 189-224, 1994.
- [4] Davidon, W.C., "Variable Metric Method for Minimization," *A.E.C. Research and Development Report*, ANL-5990, 1959.
- [5] Fletcher, R., "A New Approach to Variable Metric Algorithms," *Computer Journal*, Vol. 13, pp. 317-322, 1970.
- [6] Fletcher, R., "Practical Methods of Optimization," Vol. 1, *Unconstrained Optimization*, John Wiley and Sons, 1980.
- [7] Fletcher, R. and M.J.D. Powell, "A Rapidly Convergent Descent Method for Minimization," *Computer Journal*, Vol. 6, pp. 163-168, 1963.
- [8] Goldfarb, D., "A Family of Variable Metric Updates Derived by Variational Means," *Mathematics of Computing*, Vol. 24, pp. 23-26, 1970.
- [9] Shanno, D.F., "Conditioning of Quasi-Newton Methods for Function Minimization," *Mathematics of Computing*, Vol. 24, pp. 647-656, 1970.

See Also

[fminsearch](#) | [optimoptions](#) | [optimtool](#)